

Pipeline:

Data Preparation & Sanitization

- Initially reviewed the raw dataset and COCO-formatted JSON files containing bounding box coordinates and architectural labels.
- Developed a script to parse coordinates and crop individual objects for the ResNet classification stage
- Implemented a sanitization script to remove characters incompatible with Windows environments (e.g., colons, square brackets, spaces) from filenames and JSON paths
- Built a transformation engine to convert standard pixel bounding boxes (x, y, w, h) into YOLO-normalized center-point coordinates (cx, cy, w, h)

Preprocessing & Handling Imbalance

- Instead of standard resizing, which skews architectural shapes, I utilized **letterbox padding** to standardize image sizes while preserving original proportions.
- Used a Stratified Split (70/15/15) to ensure all five cabinet classes, including rare ones, were proportionally represented across Train, Val, and Test sets
- **For ResNet:** Applied inverse-frequency weights to the Cross-Entropy Loss function
- **For YOLO:** Implemented **Image Oversampling**, duplicating images with rare cabinet types in the training set to force the model to prioritize minority features

Architecture

- Evaluated **ResNet-18** as a baseline and **ResNet-50** for potentially better accuracy with complex features. **YOLOv8n** was selected for detection due to its small size and efficient architecture for training and deployment
- Selected **Adam** for its efficiency with sparse gradients and adaptive learning rate capabilities

What didn't work?

- Initial tests with random splits caused "blind spots" where rare classes were missing from training or evaluation.
- Simple cropping and resizing distorted the furniture geometry, making identification unreliable

Challenges

- ResNet classifiers occasionally struggled with Base vs. Wall cabinet distinction because tight crops removed the floor-level height context
- Main issues - saving weights in Kaggle, leading to potential overfitting on the final ResNet-18 inference weights.

Future work

- Using OCR to read text labels (e.g., "BC-24") alongside images to reach near-perfect accuracy

- Implementing a sliding window to process small blueprint sections at full resolution, keeping tiny grid lines sharp
- Experimenting with larger YOLO models and more data to mitigate overfitting

Results

Yolo

```

Class      Images Instances     Box(P)      R      mAP50    mAP50-95): 100% ----- 2/2 1.
4s/it 2.8s5.1s
      all       25     257   0.732   0.684   0.678   0.549
      lc_bcabo   15     154     0.72    0.87    0.73   0.562
      lc_wcabo    9      48     0.81   0.833   0.828   0.713
      lc_muscabinso   7      45   0.749   0.933   0.933   0.867
      lc_wcabcub   2      10   0.649     0.1   0.219   0.0541
Speed: 7.8ms preprocess, 12.4ms inference, 0.0ms loss, 20.9ms postprocess per image
Results saved to /kaggle/working/runs/detect/val
--- FINAL TEST METRICS ---
mAP50: 0.6775
mAP50-95: 0.5492
Precision: 0.7323
Recall: 0.6842








```


We can see that here the model, because of their fine-grained internal grid details and thin line weights, the detector often failed to distinguish them from standard wall cabinets at the current resolution. In general i think this is good results.

ResNet50

Evaluating on Unseen Test Data...

	precision	recall	f1-score	support
lc_bcabo	1.00	0.99	0.99	215
lc_bcabocub	1.00	1.00	1.00	2
lc_muscabinso	1.00	1.00	1.00	41
lc_wcabcub	0.82	1.00	0.90	14
lc_wcabo	1.00	1.00	1.00	45
accuracy			0.99	317
macro avg	0.96	1.00	0.98	317
weighted avg	0.99	0.99	0.99	317





Here we can see that the results are very good - only 3 mistakes. Also here the imbalance problem solved perfectly. However here could be problem of overfitting